

Duplicate Question Pairs Detection Report

Tapas Mandal

1 Introduction

The objective of this project is to identify whether two questions are duplicates or not. The dataset is taken from Quora. The steps include pre-processing, feature engineering, visualization, dimensionality reduction, and classification using Random Forest and XGBoost.

2 Importing Libraries and Data

We start by importing necessary libraries and loading the dataset.

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('/content/questions.csv', on_bad_lines='skip')
df = df.sample(40000)
```

Here, we import Python libraries and load 40,000 random samples from the dataset for analysis.

3 Text Preprocessing

We clean text by expanding contractions, removing special characters, and handling HTML tags.

```
from bs4 import BeautifulSoup
import re
```

```
def preprocess(q):
    q = str(q).lower().strip()
    q = q.replace('%', ' percent').replace('$', ' dollar'
    )
    q = BeautifulSoup(q, 'html.parser').get_text()
    q = re.sub(r'\W+', ' ', q).strip()
    return q

df['question1'] = df['question1'].apply(preprocess)
df['question2'] = df['question2'].apply(preprocess)
```

This ensures the text is cleaned and standardized before feature extraction.

4 Feature Engineering

We generate token-based, length-based, and fuzzy features to capture similarity between questions.

```
from nltk.corpus import stopwords
from fuzzywuzzy import fuzz
import distance
import nltk
nltk.download('stopwords')

# Example: Length features
def fetch_length_features(row):
    q1_tokens = row['question1'].split()
    q2_tokens = row['question2'].split()
    return [abs(len(q1_tokens)-len(q2_tokens)),
            (len(q1_tokens)+len(q2_tokens))/2]
```

Features like common words, fuzzy ratios, and substring lengths are extracted to represent question pairs numerically.

5 Dimensionality Reduction and Visualization

We apply t-SNE to reduce features to 2D/3D for visualization.

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.manifold import TSNE

X = MinMaxScaler().fit_transform(new_df[features])
tsne2d = TSNE(n_components=2).fit_transform(X)
```

t-SNE helps visualize how duplicate and non-duplicate questions cluster in feature space.

6 Model Training

We train Random Forest and XGBoost classifiers on the processed features.

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier

x_train, x_test, y_train, y_test = train_test_split(
    final_df.iloc[:,1:], final_df.iloc[:,0], test_size
    =0.2, random_state=42)

rf = RandomForestClassifier()
rf.fit(x_train, y_train)
print("RF Accuracy:", accuracy_score(y_test, rf.predict(
    x_test)))

xgb = XGBClassifier()
xgb.fit(x_train, y_train)
print("XGB Accuracy:", accuracy_score(y_test, xgb.predict(
    x_test)))
```

Both models are evaluated. Random Forest and XGBoost provide strong performance for classification.

7 Conclusion

The project successfully applied preprocessing, feature engineering, visualization, and machine learning models to detect duplicate questions. The models achieved good accuracy, demonstrating the effectiveness of the engineered features.